

## Generative AI Analysis Report

S. Palis

Udacity Master's Degree in AI Capstone Generative AI Applications

February 13, 2026

GitHub repository:

<https://github.com/MinervaRose/generative-ai-applications>

### Overview

This project implements a generative modeling system capable of synthesizing satellite imagery using a convolutional Variational Autoencoder (VAE). Variational autoencoders are probabilistic generative models that learn a latent representation of data by optimizing a combination of reconstruction fidelity and latent regularization (Kingma & Welling, 2014). The goal of this system is to model the distribution of Earth-observation imagery and generate visually plausible synthetic terrain textures. Such generative simulation can support robustness testing and exploratory analysis in automated sensing pipelines. The model is trained on a publicly available satellite dataset and evaluated qualitatively through reconstruction, sampling, and latent interpolation.

### Dataset Description

The dataset used in this project is the EuroSAT satellite image dataset, a publicly available collection of labeled Earth-observation images organized into land-use categories. The RGB subset contains approximately 27,000 images across ten terrain classes, including forest,

river, residential, and agricultural regions. Each image has a resolution of  $64 \times 64$  pixels, making the dataset computationally efficient while still preserving meaningful spatial structure.

The dataset is well suited for generative modeling because it exhibits consistent color distributions, texture patterns, and regional variation typical of remote sensing imagery. These properties allow a VAE to learn coarse terrain statistics without requiring extremely high-resolution inputs. The dataset is publicly accessible and appropriate for academic experimentation, supporting reproducibility and ethical reuse.

### **Model Design and Training Approach**

The system uses a convolutional  $\beta$ -Variational Autoencoder. VAEs encode input data into a latent probability distribution parameterized by a mean and variance, enabling stochastic sampling and generative reconstruction (Kingma & Welling, 2014). A 128-dimensional latent space is used to increase representational capacity relative to smaller bottlenecks, improving reconstruction fidelity while maintaining a structured latent manifold.

The training objective follows the  $\beta$ -VAE formulation, which adjusts the relative weight of the Kullback–Leibler divergence term to balance reconstruction accuracy and latent regularization (Higgins et al., 2017). A reduced KL weight ( $\beta = 0.1$ ) is applied to prioritize visual fidelity while preserving generative structure. Reconstruction error is measured using mean squared error, which is more appropriate for continuous image intensities than binary cross-entropy.

Training uses an Adam optimizer and early stopping based on validation loss to prevent overfitting. Loss curves confirm stable convergence and active use of the latent space, consistent with best practices in deep generative modeling (Goodfellow et al., 2016).

## Output Evaluation and Interpretation

Reconstructed images preserve global terrain structure and color distribution, indicating that the encoder captures meaningful spatial statistics. Large land-cover regions remain recognizable after reconstruction, although fine details such as roads and sharp boundaries are smoothed. This blur is a known characteristic of VAEs trained with pixel-level objectives, which tend to average uncertain high-frequency information.

Synthetic samples generated from random latent vectors resemble plausible satellite textures and exhibit spatial coherence without visible decoding artifacts. Latent interpolation produces smooth visual transitions, demonstrating that the learned representation is continuous rather than memorizing individual samples. These behaviors indicate successful learning of a generative manifold, even though class-specific semantic precision remains limited.

## Ethical Considerations and Responsible Use

Generative satellite imagery raises ethical concerns related to misuse, surveillance, and misrepresentation. Synthetic Earth-observation data could be used to fabricate geographic evidence or manipulate perception in political or environmental contexts. Responsible deployment of generative systems requires transparency about synthetic content and clear separation from real observational data (Floridi et al., 2018).

Additionally, generative models trained on real-world datasets may inherit biases present in the underlying data distribution (Mehrabi et al., 2021). If certain terrain types are overrepresented, the model may implicitly encode skewed assumptions about geographic prevalence. Ethical AI development requires awareness of dataset composition and explicit communication of model limitations.

## Limitations and Future Improvements

The primary limitation of this model is the loss of high-frequency detail. VAEs optimized with pixel-level reconstruction objectives favor smooth averages, reducing sharp structures and semantic precision. Increasing resolution, integrating perceptual loss functions, or hybridizing with adversarial training could improve visual fidelity.

Future work could incorporate conditional generation to produce class-aware samples, enabling more controlled terrain simulation. Higher-resolution datasets and multi-band spectral inputs may also improve realism and scientific relevance. Despite these limitations, the current system demonstrates stable generative behavior and provides a foundation for more advanced remote sensing synthesis.

## References

- Floridi, L., Cowls, J., Beltrametti, M., Chatila, R., Chazerand, P., Dignum, V., Luetge, C., Madelin, R., Pagallo, U., Rossi, F., Schafer, B., Valcke, P., & Vayena, E. (2018). AI4People—An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations. *Minds and Machines*, 28(4), 689–707.  
<https://doi.org/10.1007/s11023-018-9482-5>
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.
- Higgins, I., Matthey, L., Pal, A., Burgess, C., Glorot, X., Botvinick, M., Mohamed, S., & Lerchner, A. (2017).  $\beta$ -VAE: Learning basic visual concepts with a constrained variational framework. *Proceedings of the International Conference on Learning Representations (ICLR)*. <https://openreview.net/forum?id=Sy2fzU9gl>

Kingma, D. P., & Welling, M. (2014). Auto-encoding variational Bayes. *Proceedings of the International Conference on Learning Representations (ICLR)*.

<https://arxiv.org/abs/1312.6114>

Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on bias and fairness in machine learning. *ACM Computing Surveys*, 54(6), 1–35.

<https://doi.org/10.1145/3457607>

Rezende, D. J., Mohamed, S., & Wierstra, D. (2014). Stochastic backpropagation and approximate inference in deep generative models. *Proceedings of the International Conference on Machine Learning (ICML)*. <https://arxiv.org/abs/1401.4082>